

Using kernel methods to visualise crime data

Kieran Martin^{a,*} and Martin Ralphs^b

^a*United Kingdom Office for National Statistics, Strategy and Standards Directorate, Hampshire PO15 5RR, UK*

^b*United Kingdom Office for National Statistics, Strategy and Standards Directorate, 1 Drummond Gate, London, UK*

Abstract. Kernel smoothing methods can be used to visualise complex point event data that are distributed geographically. The visualisations generated by these methods allow users to observe trends in the data which may not be apparent from numerical summaries or by visualising the point data themselves.

Crime data in the UK has recently been made available to the public for the smallest postal delivery areas in the country (post-codes). Visualisations of the number of recorded events at each postcode for a given area are provided online. In this paper we apply some widely used kernel smoothing methods to this data, to assess the use of kernel smoothing methods as applied to open access data.

Kernel smoothing methods are applied to crime data from the Greater London metropolitan area using freely available packages for the open source R program. We also investigate the utility of using simple methods to smooth the data over time.

The kernel smoothers used are demonstrated to provide effective and useful visualisations of the data, and it is proposed that these methods should be more widely applied in exploratory data visualisation in the future.

Keywords: Data smoothing, kernel density estimation, recorded crime statistics, open source tools

1. Introduction

With the increasing level of access to data, users have an unprecedented access to statistics produced across government. For example, there are currently over 8,000 data sets available at data.gov.uk. This increased access presents issues with confidentiality; releasing such data should not compromise the identity of individuals or businesses from which statistics are gathered. This issue has been discussed in a UK government White Paper [2]. The government wishes to provide more transparent data, but avoid compromising people's identity. Large data releases are also likely to contain mistakes [6]. Approaches to releasing administrative data thus far tend to have focused on releasing data assigned to a specific location, often released in a less disclosive form by attaching multiple records to one geographical location.

Raw point estimates can be difficult to understand, especially when they are given over time or space. Simple data summaries (means over an area) can improve comprehension, but doing so can remove some of the spatial patterns that were of interest initially, as well as privilege particular geographical boundaries. One method of summarising data over geography which avoids this problem is to “smooth” the data, creating summaries over small geographical areas which blend together naturally. There have been some investigations into this by the Office for National Statistics (ONS); Carrington et al. [3] produced smoothed estimates for mortality where estimates were smoothed by simple averages, and more recently Ralphs [8] applied kernel smoothing methods to business data.

Kernel density estimation is a so called “non-parametric” method for summarising data gathered over one or more dimensions. It can be used in multiple contexts, such as estimating the shape of a density function, or looking for underlying trends in data. Kernel density estimation is a descriptive measure, applying to observed data, and does not necessarily provide accurate predictions for future data as parametric methods can.

*Corresponding author: Kieran Martin, Strategy and Standards Directorate, Room 2400, Office for National Statistics, Segensworth Road, Titchfield, Fareham, Hampshire PO15 5RR, UK. Tel.: +44 1328 44 760; E-mail: kieran.martin@ons.gov.uk.

One use of kernel density estimation can be to visualise patterns in spatial data, such as levels of unemployment or poverty across the country. In this report we illustrate the use of kernel smoothing methods by application to a publically available data set. Smoothed data will naturally reduce how easy it is to identify individuals, and help control against mistakes in releases by averaging out errors.

The British police keep a record of all reported crime in each constabulary. This is one method used in the United Kingdom for recording crime levels; the British Crime Survey is also used to supplement recorded crime by surveying people's experience of crime in the last 12 months.

For the last two years, police.uk has released monthly information on the number of criminal events that have been recorded by police authorities in England and Wales. This release combines all crimes recorded across that month, and the events are then associated with the centre of the postcode in which said crime was reported to have occurred. While this is a useful visualisation at a local level, as of yet there is no real summary of crime statistics across an area, so that like with like comparisons are difficult. Comparing one month with another in a particular area can only be done by examining the raw totals of crimes recorded for an area.

This paper describes the application of kernel smoothing methods to crime data. In Section 2 the general theory of the kernel density methods used is briefly described, and how the methods were applied to the crime data is outlined in Section 3. Section 4 presents results for crime across the greater London metropolitan area for several crime types, as well as for a smaller area, in Camden. Section 5 presents results when the crime data has been averaged over time, to try to dampen the effect of small spikes in crime in a particular month. Section 6 details conclusions and proposals for further investigation.

2. Kernel density estimation for spatial point process patterns

Kernel density estimation involves applying a function (known as a “kernel”) to each data point, which averages the location of that point with respect to the location of other data points. The method has been extended to multiple dimensions and provides an effective method for estimating geographical density. There is a good deal of literature discussing kernel density

estimation, such as Simonoff [11]; kernel density estimation was applied in this paper using the package *spatstat* in R [1].

Figure 1 provides an illustration of how this works in practice. The figure shows a number of data points scattered across a two dimensional space. In kernel density estimation, a smooth, curved surface in the form of a kernel function (which looks a little like a tent in the diagram) is placed at each data point. The value of the kernel function is highest at the location of the point, and diminishes with increasing distance from the point.

Now follows a brief introduction to the mathematical underpinnings of kernel estimation. Given a data set $(\mathbf{x}_1, \dots, \mathbf{x}_n)^T$ where $\mathbf{x}_i = (x_{i1}, \dots, x_{id})^T$, d being the dimension being averaged over. Kernel density estimation uses a pre-defined kernel, \mathbf{K} , and a bandwidth matrix \mathbf{H} to produce an estimate of the density of said data. The estimate of the density of the data at a point \mathbf{x} is then:

$$\hat{F}(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \mathbf{K}_{\mathbf{H}}(\mathbf{x} - \mathbf{x}_i),$$

where,

$$K_{\mathbf{H}}(x) = |\mathbf{H}|^{-1/2} K(\mathbf{H}^{-1/2} x).$$

This is effectively a weighted average over an area determined by the bandwidth matrix, \mathbf{H} , where observations closer to the centre point are weighted more highly. How much individual observations are weighted depends on the choice of the kernel. Popular choices include the Epanechnikov kernel and the Gaussian kernel, which is used in this report. The Gaussian kernel was chosen as it is commonly used and is the kernel used in *spatstat*. It has the following form:

$$K(x) = (2\pi)^{-d/2} \exp\left(-\frac{1}{2}x^T x\right).$$

This kernel weights observations by their Euclidian distance from \mathbf{x} , with the weight reducing at an exponential rate as the distance increases. Figure 2 gives a visualisation of the kernel, where the height indicates the amount of weight given to an observation at that location.

The data used in this report can be considered to be a spatial point process. This is a data set distributed in 2 dimensions (or three, with time), that has an additional “intensity” attribute at each observation. For crime data, this intensity is the number of incidents observed at that point in space: other examples might be the recorded heights of trees at particular locations.




Fig. 1. Visualisation of kernel smoothing process.

The intensity of the crime can be smoothed over space using kernel methods, by calculating weighted kernel estimates via the following formula:

$$\frac{\sum_i m_i \mathbf{K}_H(\mathbf{x} - \mathbf{x}_i)}{\sum_i \mathbf{K}_H(\mathbf{x} - \mathbf{x}_i)},$$

where m_i is the intensity associated with \mathbf{x}_i . This is known as the Nadaraya-Watson estimator [7].

Bandwidth selection can be performed manually by examining the output and selecting an appropriate value, or by using various automated methods [10]. Investigation into possible automated methods, such as Diggle's method [4] found that the bandwidths obtained over smoothed the results at small geographies, and were computationally difficult to obtain at high geographies. It was found that Scott's rule of thumb [9] provided good values for the bandwidth and was easy to quickly calculate. Scott's rule has the form:

$$h_j = n^{-1/(d+4)} \sigma_j,$$

where h_j is the bandwidth for the j_{th} dimension, $j = 1, \dots, d$ ($d = 2$ in this paper), n the number of observations and σ_j is the standard deviation of the observations in the j_{th} dimension. This is a simple method, which is based on providing the optimal bandwidth for a reference distribution, in this case a multivariate normal distribution.

3. Applying kernel density estimation to crime data

The crime data as released by the home office is given in the form of the number of events associated with a particular location. While in practice these events will have not occurred in the exact same geographical location, to avoid disclosure the outputs are given in this form. This means that the data is treated as a spatial point process, with the number of events being the intensity, m .

Treating the data as a spatial process implies that each point is an observation of intensity. In this case,




Fig. 2. Visualisation of the Gaussian kernel.

the model needs to take account of the fact that most locations do not observe crimes in any given month, otherwise the smoothing process will assume that the minimum value of crimes observed at any location is 1. To account for this, all crime data over the two year period available was gathered to populate London with possible locations where crimes might occur and each location was assigned an intensity of 0; then, in any given month, those locations which recorded crimes had their intensity value replaced with the number of recorded crimes. Note that there may be some locations that in the two year period that data is available for did not report any criminal events, and thus are missed by this method, but as recorded crimes include anti-social behaviour it is hoped that the majority of postcodes are included by this method.

Two different areas were considered: the greater metropolitan London area, and a small area which includes south Camden and Euston. This latter area was chosen to demonstrate the kernel smoothing method applied to smaller geographies.

For each geographical area multiple crime types could be considered. This paper focuses on vehicle crime, but this method can easily be applied to any crime type. The chosen bandwidth depends on the total number (and spread) of points smoothed over, so is constant across crime type, but varies depending on the level of geography considered.

4. Kernel smoothing crime over London

Crime maps were generated for each month from 03/2011–08/2011 for vehicle crime. These are dis-

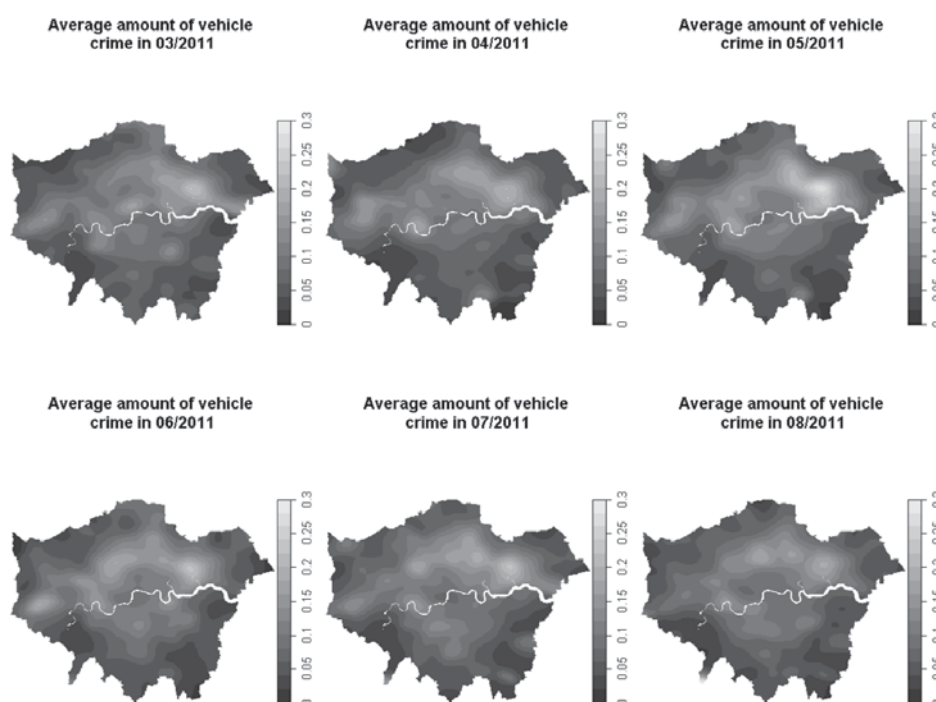


Fig. 3. Vehicle crime in London between 03/2011 and 08/2011.

played for London in Fig. 3. The scale gives the average number of events in that area.

This figure demonstrates the power of this technique. It can be seen that while there is some variance in crime levels between different months, the hot spots remain in similar locations. In particular, vehicle crime remains high in north east London throughout the six month period considered. Visualisations at any level which include greater than 10,000 crimes are not currently available at police.uk.

Figure 4 displays the results for vehicle crime in Camden. The figure also includes the recorded crime events for that month, represented by circles, which are sized proportional to the number of recorded events at the centre of the circle. This allows a comparison between currently available crime visualisations and the new method.

It is clear that over small areas crime is more variable, which is possibly unsurprising, as smaller areas will show small changes each month, while across London these minor changes will be averaged out. In particular Fig. 4 shows a large spike in vehicle crime in March (25 events occurred on one road), which is reflected in the smoothing method.

The contours act as an informative summary of the point data, providing more information on the areas in which most events are occurring.

5. Results smoothed over time

Kernel estimation methods could hypothetically be used to smooth over time as well as over space, but those methods are not currently available for spatial point processes in R. Instead, we investigated averaging the recorded crime events over time and then applying kernel methods to the averaged data.

Initially a simple average of the amount of crime in each location over the 6 month period 03/2011–08/2011 was calculated and then kernel smoothing methods were applied to the averaged values for both London and Camden, plotted in Figs 5 and 6.

These figures contain the trends that seemed to be occurring in the previous time series but also display trends which are less obvious at a month to month level. Figure 6 demonstrates that there is an area in South East Camden where vehicle crimes are occurring more frequently; this is less apparent when observing individual months.

Another possible option for averaging data over time is to use a centralised moving average. For a given lag (two was used here), a centralised moving average gives the value for each month as the average of the values recorded for that month, and the two months before and after (with the number depending on the lag).

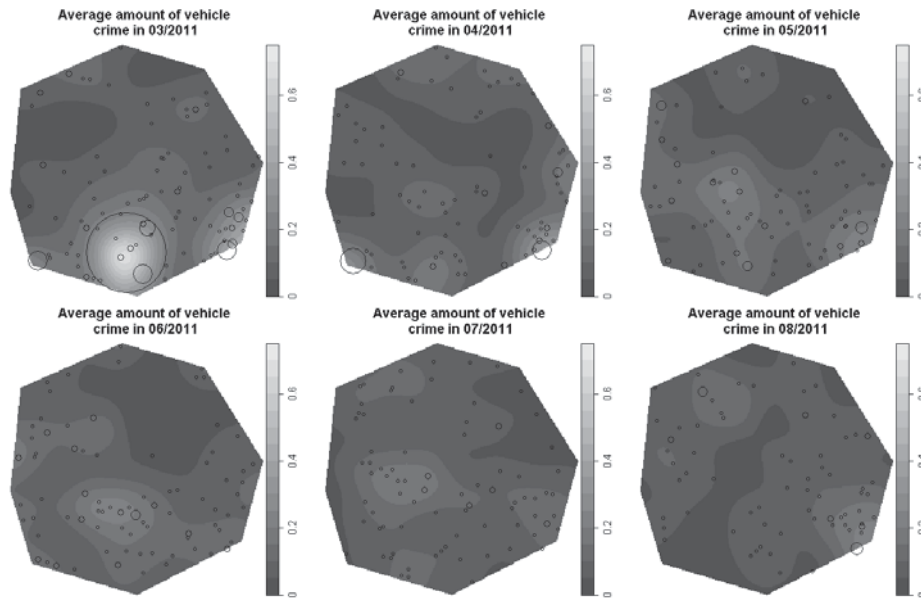


Fig. 4. Vehicle crime in Camden 03/11-08/11.

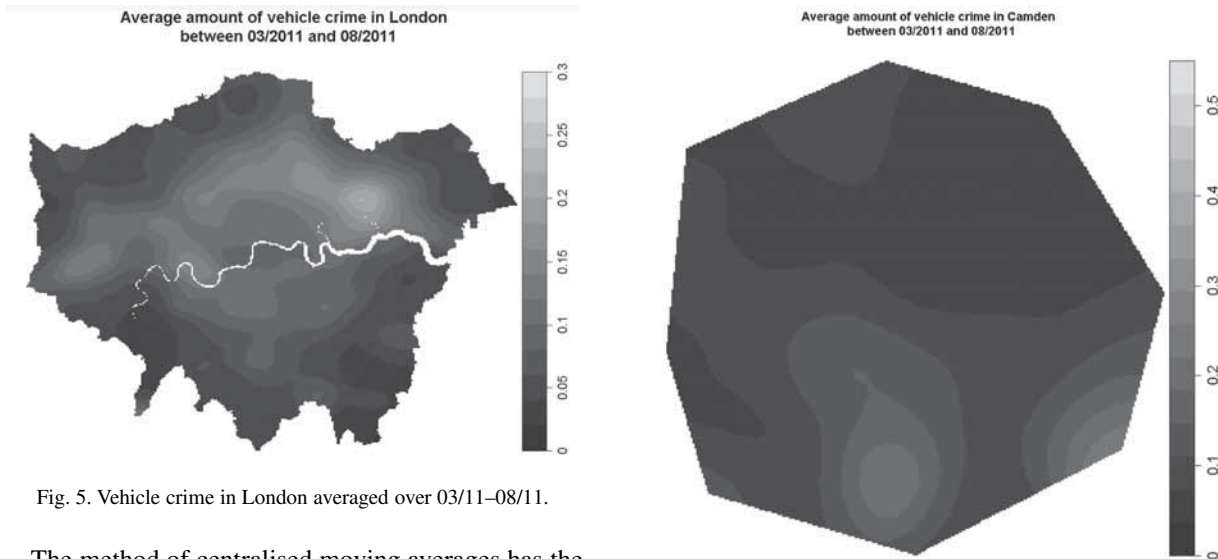


Fig. 5. Vehicle crime in London averaged over 03/11-08/11.

Fig. 6. Vehicle crime in Camden averaged over 03/11-08/11.

The method of centralised moving averages has the advantage of smoothing out individual spikes in each month, and hopefully makes any underlying trend easier to distinguish. Different weightings can be applied to each month, and the lag can be varied depending on the practitioner's understanding of how long it might take for a trend to appear. One downside of this method is that, as it requires the months before and after the current one, for a complete time series the last and first two (depending on the size of the lag) observations will not be using all data, and should not be displayed. In this report we have chosen a subset of the available data so this does not affect the plots we display, but it

is a practical consideration that should be kept in mind when this technique is applied.

Figures 7 and 8 display the centralised moving averaged crime data, smoothed over London and Camden.

Figure 7 shows less month to month variation, while allowing for gradual change, and is possibly a better visualisation of crime data than Figs 5 and 3 for appreciating the trends over time of crime in London. The figure is a compromise between simply averaging over the six months and not averaging at all. The

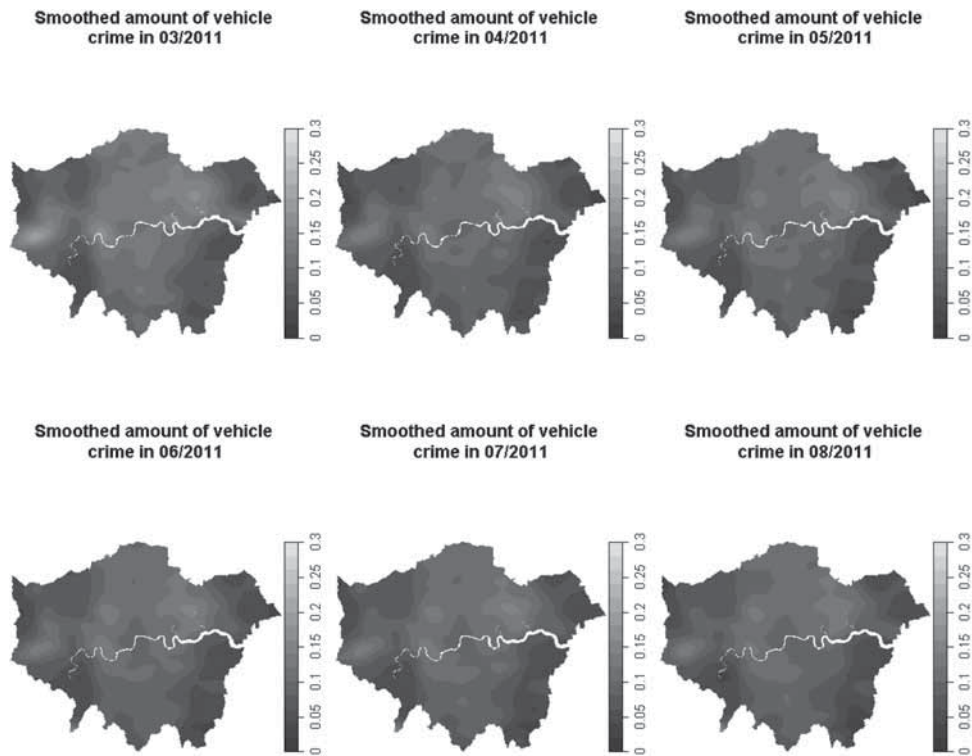


Fig. 7. Moving average of vehicle crime in London from 03/11–08/11.

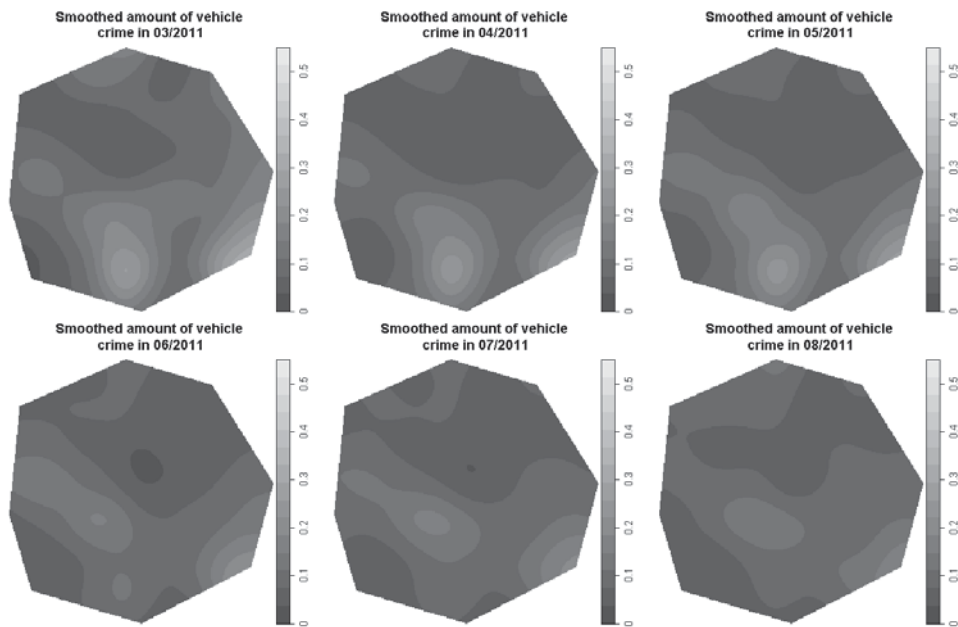


Fig. 8. Moving average of vehicle crime in Camden from 03/11–08/11.

6 month average removes any short term trends in the data, while the simple month by month summary can be too sensitive to mild variation.

The downside of this method is that outliers may have too strong an effect. Figure 8, which displays vehicle crime in Camden, does show less noise than Fig. 4, but picks out a trend of vehicle crime being concentrated in the south for the first three months which is due to the extreme spike in vehicle crime in 03/2011 rather than a consistent pattern over that period of time. The same problem affects Fig. 6, which has the same hot spot.

6. Conclusions and future research

This investigation has demonstrated that kernel smoothing methods can be a powerful tool for visualising and understanding patterns in crime data. This method could be applied to additional data sets, such as mortality data and business data; both data sets have had smoothing techniques applied them in the past. Potentially any data set with a geographical element could have kernel smoothing methods applied to it.

The method allows users to compare the crime both over time and within an area; this provides advantages when looking for patterns which may be less apparent in point data. It also potentially aids in avoiding disclosure issues; if data is released only in smoothed format then it might be more difficult to identify individuals.

Kernel estimates are relatively easy to calculate in R, which is freely available. Interpreting patterns and differences between areas can be much more difficult when only point estimates are available. This method has been put to exploratory use here, but it can potentially be used in an analytic fashion [5]. Further investigation could also be put into smoothing data over time: the methods used here are somewhat ad hoc, and theoretically kernel methods could be applied to time as well as space.

Additional visualisations could be used to overlay the kernel smoothing maps over actual geography, such as Google maps, allowing users to see the crime levels overlaid on their areas. This will enable users to look for explanations of why an area has a low or high level of vehicle crime: a car park or busy road might have more vehicle crime associated with it, for example.

References

- [1] A. Baddeley and R. Turner, Spatstat: An R package for analyzing spatial point patterns, *Journal of Statistical Software* **12** (2005), 1–42.
- [2] Cabinet Office, Open data white paper: unleashing the potential, HMSO, London, 2012, (available at <https://www.gov.uk/government/publications/open-data-white-paper-unleashing-the-potential>).
- [3] A. Carrington, P. Heady, M. Ralphs, N. Rahman and B. Mitchell, Smoothing of standardised mortality ratios, National Statistics Methodological Series No 35, Office for National Statistics, United Kingdom, 2007.
- [4] P.J. Diggle, A kernel method for smoothing point process data, *Applied Statistics (Journal of the Royal Statistical Society series C)* **34** (1985), 138–147.
- [5] P.J. Diggle, *Statistical analysis of spatial point patterns*, Hodder Arnold, London, 2003, second edition.
- [6] D. Hand, Open data: problems and promise, *Methods News. Newsletter for the ESRC National Centre for Research Methods*, Summer 2012.
- [7] E.A. Nadaraya, On estimating regression, *Theory of Probability and its Applications* **9**(1) (1964), 141–142.
- [8] M. Ralphs, Using kernel density estimation of inter-departmental business register data to explore functional economic geographies in Greater London, Proceedings of the Sixteenth Government Statistical Service Methodology Symposium, 2011, (available at <http://www.ons.gov.uk/ons/about-ons/get-involved/events/events/sixteenth-gss-methodology-symposium-6-july-2011/sixteenth-gss-methodology-symposium-papers-and-presentations/index.html>).
- [9] D.W. Scott, Multivariate density estimation. Theory, practice and visualization, Wiley, New York, 1992, p. 152.
- [10] S.J. Sheather, The performance of six popular bandwidth selection methods on some real data sets (with discussion), *Computational Statistics* **7** (1992), 225–250, 271–281.
- [11] J.S. Simonoff, Smoothing methods in statistics, Springer, 1996.

Copyright of Statistical Journal of the IAOS is the property of IOS Press and its content may not be copied or emailed to multiple sites or posted to a listserv without the copyright holder's express written permission. However, users may print, download, or email articles for individual use.